

Pre-registration (DRAFT) — An Anomaly-Residual Search for Unexplained Transit and Photometric Signatures Around Old, Quiet K Dwarfs

Status: DRAFT / in preparation — NOT yet registered. Phase 2 of the program whose Phase 1 (white dwarfs) is registered at OSF [10.17605/OSF.IO/6YH7R](https://doi.org/10.17605/OSF.IO/6YH7R). This plan is to be frozen and registered on OSF **before any K-dwarf light curve is analysed**; until then it is edited freely in this repository. It reuses the *validated* population-agnostic pipeline core (see `pipeline/core/`, regression-tested in `pipeline/runners/validate_wd.py`); the K-dwarf-specific parts are a new population plugin, not a fork.

0. Summary

A pre-registered, mechanism-agnostic search for anomalous transit and photometric signatures around **old, quiet main-sequence K dwarfs** — the longest-lived, most thermally-stable, and photometrically-quietest *living* stellar hosts. Phase 1 applied this anomaly-residual methodology to white-dwarf remnants and returned a clean, fully-explained null with a quantitative upper limit. Phase 2 carries the **same validated machinery** to a living host, with **transit morphology as the primary channel**. We search for departures from the natural model — most powerfully, transit light curves whose *shape* no natural occulter (planet, eclipsing binary, disintegrating body, starspot) can produce — and try to explain each away through a fixed natural-explanation battery, with detection thresholds set by an empirical null and injection-recovery, never by inspecting candidates.

1. The search target and its assumptions

- **The target is the observable anomaly** — an unexplained transit/photometric departure — not a presupposed structure, mechanism, or agent (as in Phase 1).
- **The agnostic stance, extended.** We are agnostic about: (a) *mechanism* — no assumed engineering; (b) *engineering vs. mere existence* — we detect departures in the star’s light and do not presuppose they are built; (c) *timescale* — we assume no schedule of activity. A consequence we state plainly: over our ~decade observational baselines, any process on

decade timescales is effectively **static**, so the realistic primary signal is a *steady* anomaly (a persistent anomalous transit shape), with genuine *variability* (a changing transit depth or morphology) the highest-value-but-lowest-prior bonus, not the expectation.

- **The one prior we retain: time for life to have originated.** This is not a mechanism or timescale assumption — it is the precondition for there being anything to find at all. K dwarfs uniquely maximise it: 15–40 Gyr main-sequence lifetimes mean every K dwarf has had ample time and none has evolved off the main sequence. The sample is further weighted toward *old* K dwarfs (§3).

- **What we refuse to assume.** We do not assume life requires planets, photosynthesis, visible light, or any particular biochemistry. None of these enters the detection — which measures only departures in the star’s integrated light — so the search takes no position on them.
- **Why K dwarfs, and the honest limitation.** Among *living* hosts, K dwarfs offer the longest stable lifetime, the quietest baseline (calmer than flaring M dwarfs), and — being small — the deepest, most morphologically-resolvable transits. The cost, stated up front: a living star’s cool photosphere is bright in the infrared, so the IR-excess channel that anchored Phase 1 loses its contrast. **Phase 2 is therefore transit-anchored, not excess-anchored.** We further acknowledge the transit channel is comparatively well-trodden; our differentiator is not the channel but the *discipline* — calibrated empirical-null thresholds and a mandatory, pre-registered explain-away battery applied uniformly.

2. Research questions

- **RQ1 (morphology).** Do any K dwarfs show a periodic transit whose *shape* — asymmetry, flat-bottomedness, box-versus-U, anomalous duration or secondary structure — cannot be reproduced by any natural occulter (sphere, ring, eclipsing binary, disintegrating body) after those are explicitly fitted and excluded?
- **RQ2 (photometric departure).** Do any show a persistent or time-varying photometric signature departing from the natural rotation/granulation/activity model that survives the battery?
- **RQ3 (upper limit).** With no unexplained anomaly, what upper limit f_{max} can we place on the prevalence of anomalous occulters/departures around old, quiet K dwarfs, as a function of depth and period?

3. Sample and inclusion

- **Parent sample.** Main-sequence K dwarfs from Gaia DR3: T_{eff} 3900–5300 K, $\log g > 4.3$ (dwarf, excluding subgiants/giants), with clean astrometry ($\text{RUWE} < 1.4$, $\text{parallax_over_error} > 10$), plus the equivalent absolute-G versus ($G_{\text{BP}} - G_{\text{RP}}$) main-sequence box. *Scale (this repository, sized pre-data):* ~12,600 at $G < 11$; ~22,000 within 100 pc; ~178,000 at $G < 13$. Target finding is not the binding constraint. The exact constraints (the colour-magnitude box, the RUWE and parallax cuts, and the hard magnitude/volume limit) are committed as an **immutable, checksummed manifest** of `source_ids` before any light curve is pulled — the frozen sample against which all later results are reproducible, identical in spirit to the Phase-1 `wd_sample` manifest.
- **Data-quality / channel requirement.** A usable TESS and/or Kepler light curve (the transit instrument). The Kepler-field subset (4-year continuous photometry) is the morphology gold standard and is treated as a distinct, deepest tier.
- **The goal is the whole population; staging is a resource constraint, not a scientific one.** We impose no scientific exclusion beyond the minimal *youth floor* below and the data-quality requirement above. Which stars have been *analysed* at any moment is set by available compute; the asymptotic target is the entire K-dwarf census.

- **Youth floor (the one real cut — soft and pre-specified).** We exclude only genuinely young, active stars, for two reasons that point the same way: the *life prior* (a star must have *existed* long enough for life to have originated, regardless of how long it will ultimately live) and *data quality* (young K dwarfs are magnetically active — the worst photometric noise). Applied in two pre-specified stages: a coarse Gaia-only proxy in the manifest (kinematics; Gaia photometric-variability amplitude), refined before the transit search by rotation-period gyrochronology and activity indicators (GALEX, ROSAT/eROSITA, Ca II) where catalogued. Thresholds are fixed from the parent-sample distributions, never tuned to candidates.
- **No hard “quiet” cut — noisy stars self-weight.** Rather than excluding active stars, we include every star above the youth floor and let the per-star injection-recovery completeness C_i (§5) do the work: a transit is harder to recover in an active light curve, so a noisy star earns a low C_i and contributes little to $f_{\max}$ automatically (the Phase-1 self-weighting principle — under-probed objects weight toward zero). The limit stays honest over whatever has been analysed.
- **Resource-triggered, nested, staged analysis order (so iteration stays confirmatory).** Because the goal is everything and the binding constraint is compute, we pre-register a *nested* sequence of tiers (cleanest/brightest first: $T_1 \subset T_2 \subset \dots \subset$ the full sample), analysed in that fixed order, with the expansion trigger being *available resources*, declared in advance — never the result of a prior tier. The survey-wide trial-factor / family-wise threshold (§5) is recomputed for each expanded sample (alpha-spending across stages). Growing the sample is therefore a pre-registered confirmatory extension (cf. the Phase-1 variability expansion), not a post-hoc tweak; what is forbidden is loosening any boundary *in reaction to a result* or to rescue a candidate.
- No assumption that an anomaly orbits in a habitable zone; all periods the data permit are searched.
- **Stated selection biases:** brightness-limited (toward nearby), the TESS/Kepler footprint, and the youth floor. $f_{\max}$ is interpreted as a prevalence among the *effectively probed, solar-neighbourhood* K-dwarf population, self-weighted by per-star completeness.

4. Channels

- **Channel B — transit morphology (PRIMARY, CALIBRATED).** BLS/TLS detection; morphology metrics (asymmetry, flat-bottom fraction, box-vs-U, duty cycle); and a **mandatory automated difference-image centroid gate** (§5). The Phase-1 machinery transfers directly: K dwarfs are point sources with clean Gaia astrometry, and a real occulter gives a *deep* transit, so the “deep $\Rightarrow$ on-target real; shallow-and-offset $\Rightarrow$ background blend, confirmed by centroiding” logic applies (the same code validated for white dwarfs, [pipeline/analysis/07–09](#)). Crucially, the sample size — tens of thousands versus Phase 1’s 157 — **graduates this from a secondary, candidate-generating channel into a primary, calibrated one:** it now supports a rigorous population-level upper limit on structural transit anomalies (RQ3, §5), the transit analogue of Phase 1’s IR-excess $f_{\max}$.

- **Channel P — photometric departure (SECONDARY).** Departures from the natural rotation/granulation/activity model; correlated/structured variability; same empirical-null calibration as Phase 1.
- **Channel A — infrared excess (CORROBORATING; extreme-outlier only, no standalone limit).** A K dwarf’s cool IR photosphere suppresses excess contrast, and debris disks are a common, *expected* natural explanation — so this channel is toothless for *marginal* signals and carries **no calibrated f_{max}** . Following the Phase-1 Channel-C precedent (an ordinal corroborating flag with no standalone threshold), it is registered as a **high-bar extreme-outlier flag**: the empirical null, calibrated against the disk-rich bulk, flags only the *far tail* — an excess outside any natural disk/companion regime, an anomalous SED shape, or (highest value) a *fluctuating* excess that a static disk cannot fake. Every flag is cross-checked against SIMBAD / the literature / known debris-disk catalogues and the already-explained ones killed (prior knowledge used as a *filter*, never as an assumption that the tail is empty). A surviving extreme excess **elevates an object only when it coincides with a Channel-B transit anomaly on the same star** — that coincidence (an otherwise-inexplicable IR excess *and* an anomalous transit on one K dwarf) is the high-value residual, far stronger than either channel alone.

5. Natural-explanation battery, statistics, and stopping rule

A flagged candidate becomes a residual only if it survives every applicable test. The order is deliberate: the cheapest, highest-yield contaminant filter runs *first and automatically*, because at this sample size nothing can be vetted by hand.

1. **Background eclipsing binary / aperture blend — automated centroid gate (FIRST, mandatory).** TESS’s $\sim 21''$ pixels guarantee that *hundreds* of off-target eclipsing binaries blend into K-dwarf apertures — the dominant transit false positive at scale. Difference-image centroiding (the flux-weighted centroid of the out-of-transit minus in-transit image) is run **automatically as the first gate**: a candidate whose dip centroid sits more than a pre-set tolerance (≈ 1 TESS pixel) from the target is killed with no human inspection. In Phase 1 this was a manual check on 3 candidates; here it is a mandatory, automated guillotine at the front of the battery.
 2. **Eclipsing binary (on-target)** — fit EB models; check for a secondary eclipse, odd-even depth differences, ellipsoidal/reflection modulation; radial velocities where available.
 3. **Stellar activity** — starspot rotational modulation, flares, faculae; cross-check against the star’s rotation period and activity indicators.
 4. **Bona-fide planet** — a real planet transit is *natural*; we flag only morphologies inconsistent with a spherical or ringed planet (noting the ringed-planet ambiguity of Arnold 2005).
 5. **Disintegrating planetesimal / dust tail** — tested by recurrence, depth evolution, and an explicit asymmetric dust-tail template (cf. WD 1145+017, KIC 12557548).
 6. **Instrumental / systematic** — TESS/Kepler systematics, momentum dumps, scattered light, aperture contamination.
- **Detection threshold and the trial factor.** With tens of thousands of stars $\times \sim 10^6$ BLS period/phase trials each, a 3-4 σ bar would bury the residual list in noise. We

therefore compose two corrections, both pre-specified as *procedures* (the resulting numbers come from the data via those procedures, never from inspecting candidates): **(a) an empirical null on the per-star BLS detection statistic, computed within noise/activity cohorts** (binned by each light curve’s photometric scatter / CDPP / activity level) rather than globally — so a quiet star is judged against other quiet stars and the active stars’ heavy-tailed noise (spots, flares, rotation) does **not** inflate the threshold for the clean ones. This per-cohort grading (the lesson from the v2 W2-offset handling, generalised) is what lets us keep every star above the youth floor without a hard activity cut: within each cohort, genomic-control inflation λ rescales for the non-Gaussian reality of TESS/Kepler systematics — as the WISE-excess null did ($\lambda \approx 10.6$ in Phase 1). And **(b) a survey-wide family-wise / look-elsewhere correction** (Gross & Vitells 2010; FDR, Benjamini- Hochberg / Storey) such that the expected number of pure-noise false alarms across the *entire frozen sample* is < 1 . The resulting threshold is expected to land in the $\approx 6\text{--}7\sigma$ regime (cf. the Kepler 7.1σ bar, Jenkins et al. 2002, for $\sim 150\text{k}$ stars) — but it is **derived from our own injection-recovery on the frozen sample, not borrowed from Kepler**, since it depends on our sample size, cadence, and period range. It is frozen before the real candidate tail is unblinded.

- **Upper limit (Channel B, now calibrated).** The Poisson zero-detection bound $f_{\text{max}}(\text{depth}, \text{period}) = 3 / \sum_i C_i$, with per-star completeness $C_i(\text{depth}, \text{period})$ measured by **injection-recovery of synthetic anomalous transits into the real light curves** — the same logic as Phase 1’s IR-excess limit, now yielding a population constraint on structural transit anomalies as a function of depth and period.
- **Stopping rule (unchanged from Phase 1):** “unexplained” means “survives the registered battery at procedure-frozen thresholds,” nothing more. We report the *full ranked residual list*, and the standing interpretation is that a residual is a target for conventional astrophysical follow-up — never, in this document or its outputs, a claim of detected intelligence.

6. Analysis plan

1. Freeze the sample $\rightarrow$ the immutable checksummed manifest (§3).
2. Pull TESS/Kepler light curves (deterministic recipe); detrend.
3. **Calibrate and freeze the thresholds before unblinding** — run the empirical-null calibration on the bulk BLS-statistic distribution (the noise floor, not the candidate tail) and the injection-recovery on the frozen sample to fix the family-wise-controlled detection threshold and the per-star completeness C_i . *This step touches only the noise floor and synthetic injections — never the real candidate tail.*
4. **Unblind:** run the BLS/TLS search, apply the frozen threshold, compute morphology metrics, and pass survivors through the natural-explanation battery (automated centroid gate first).
5. Report the full ranked residual list and the calibrated $f_{\text{max}}(\text{depth}, \text{period})$.

Same recipe-in-repo, checksum, gitignored-bulk discipline as Phase 1. The split between steps 3 (noise-floor/synthetic calibration, thresholds frozen) and 4 (real

candidates unblinded against the frozen thresholds) is the integrity crux at this sample size: it is what prevents a strange real signal from tempting a threshold tweak.

7. Outcomes and interpretation

- **Expected:** a clean, explained null plus a prevalence upper limit, with a byproduct catalogue of K-dwarf transit candidates and activity characterisation of value independent of the technosignature framing.
- **A surviving residual** is, in order of probability, a *previously-unmodelled natural* occulter or variability class — published as such, with all natural tests shown. A *fluctuating or structured* anomalous morphology would be the single highest-value outcome. No claim of artificiality is made on the basis of this pipeline.

8. Open science and reproducibility

Public repository; pinned data recipe, query text, frozen manifests, and checksums; bulk data fetched on demand (gitignored). The analysis runs on the **validated population-agnostic core** plus a new **K-dwarf plugin** — Phase 1 reused, not forked, with the WD regression test as a standing guarantee that the shared core is unchanged. AI-collaboration transcripts archived. Registered on OSF before any K-dwarf light curve is analysed; post-registration changes are dated, public amendments.

Statement of Provenance and Acknowledgments

Human accountability. Authored and directed by the sole investigator (T. Loewald), who bears full responsibility for its contents, methodological selections, and errors. Special thanks are owed to a cross-disciplinary collaborator whose insight — bringing the population-level statistical standards of epidemiology and genome-wide association studies to bear on an astrophysical problem — motivated the empirical-null and genomic-control framework this search inherits from Phase 1.

AI intellectual provenance. Two AI systems (Google Gemini, Anthropic Claude) functioned as active computational logic engines and co-designers in methodology, implementation, and adversarial review; they are credited as such while the investigator retains all accountability. The complete, lightly-redacted working transcripts are archived in the public repository.

References

(Shared methodology references as Phase 1, plus transit-specific.) - Arnold, L. F. A. 2005, *ApJ*, 627, 534 — artificial/ringed transit signatures. - Kovács, G., Zucker, S., & Mazeh, T. 2002, *A&A*, 391, 369 — Box Least Squares. - Hippke, M., & Heller, R. 2019, *A&A*, 623, A39 — Transit Least Squares. - Boyajian, T. S., et al. 2016, *MNRAS*, 457, 3988 — KIC 8462852 (anomalous, structured transits). - Vanderburg, A., et al. 2015, *Nature*, 526, 546 — disintegrating planetesimals at WD 1145+017. - Borucki, W. J., et al. 2010, *Science*, 327, 977 — Kepler. - Ricker, G. R., et al. 2015, *JATIS*, 1, 014003 — TESS. - Jenkins, J. M., Caldwell, D. A., & Borucki, W. J. 2002, *ApJ*, 564, 495 — transit detection threshold. - Efron, B. 2004, *JASA*, 99, 96 — empirical null. · Devlin & Roeder

1999, *Biometrics*, 55, 997 — genomic control. - Gentile Fusillo, N. P., et al. 2021, *MNRAS*, 508, 3877 / Gaia Collaboration — sample provenance.